

# Simple and Efficient Learning with Dynamic Neural Networks

Graham Neubig



**Carnegie Mellon University**

Language Technologies Institute




Code Examples:

<https://github.com/neubig/lxmls-2017>

# Neural Networks for Language

- Neural networks give us new tools to process data: images, speech, text
- Particularly for text, we would like to use networks with complicated structure
- And we want to go from idea to code quickly

# Example Task: Sentiment



# What is Necessary for Neural Network Training

- **define** computation
- **add** data
- calculate result (**forward**)
- calculate gradients (**backward**)
- **update** parameters

# Paradigm 1: Static Graphs (Tensorflow, Theano)

- **define**
- for each data point:
  - **add data**
  - **forward**
  - **backward**
  - **update**

# Advantages/Disadvantages of Static Graphs

- **Advantages:**
  - Can be optimized at definition time
  - Easy to feed data to GPUs, etc., via data iterators
- **Disadvantages:**
  - Difficult to implement nets with varying structure (trees, graphs, flow control)
  - Need to learn big API that implements flow control in the “graph” language

# Paradigm 2: Dynamic Graphs (Chainer, DyNet, PyTorch)

- for each data point:
  - **define**
  - **add data/forward**
  - **backward**
  - **update**

# Advantages/Disadvantages of Dynamic Graphs

- **Advantages:**
  - API is closer to standard Python/C++
  - Easy to implement nets with varying structure
- **Disadvantages:**
  - Harder to optimize graphs (but still possible, see end of presentation!)
  - Harder to schedule of data transfer, etc.




# DyNet

<https://github.com/clab/dynet>

- Dynamic graph toolkit implemented in C++ **usable from C++, Python, Scala/Java (soon Haskell?)**
- **Very fast on CPU** (good for prototyping NLP apps!), similar support to other toolkits for GPU
- Support for **easy implementation of mini-batching**, even in difficult situations

# Programming Examples

# Bag of Words (BOW)



# At Beginning of Training

```
# Start DyNet and define trainer  
model = dy.Model()  
trainer = dy.AdamTrainer(model)
```

```
# Define the model  
W_sm = model.add_lookup_parameters((nwords, ntags))  
b_sm = model.add_parameters((ntags))
```

## Trainer

Our strategy for training the model (here Adam)

## Regular Parameters

A parameter vector/matrix/tensor (here `b_sm` is size `ntags`)

## Lookup Parameters

One vector for each word (here `W_sm` has `nwords` words, vector of size `ntags`)

# Calculating the Network

```
# A function to calculate scores for one sentence  
def calc_scores(words):  
    # Create a computation graph, and add parameters  
    dy.renew_cg()  
    b_sm_exp = dy.parameter(b_sm)  
    # Take the sum of all the embedding vectors for each word  
    score = dy.esum([dy.lookup(W_sm, x) for x in words])  
    # Add the bias vector and return  
    return score + b_sm_exp
```

# Training Time

```
# Perform training over the entire corpus
train_loss = 0.0
for words, tag in train:
    # Calculate the scores for each candidate
    my_scores = calc_scores(words)
    # Cross-entropy loss function for the correct tag. my_loss is a
    # DyNet expression (we have not performed calculation yet)
    my_loss = dy.pickneglogsoftmax(my_words, tag)
    # Call the ".value()" function to perform actual calculation
    train_loss += my_loss.value()
    # Perform backward calculation and update
    my_loss.backward()
    trainer.update()
# Print the values
print("iter %r: train loss/sent=%.4f" % (ITER, train_loss/len(train)))
```

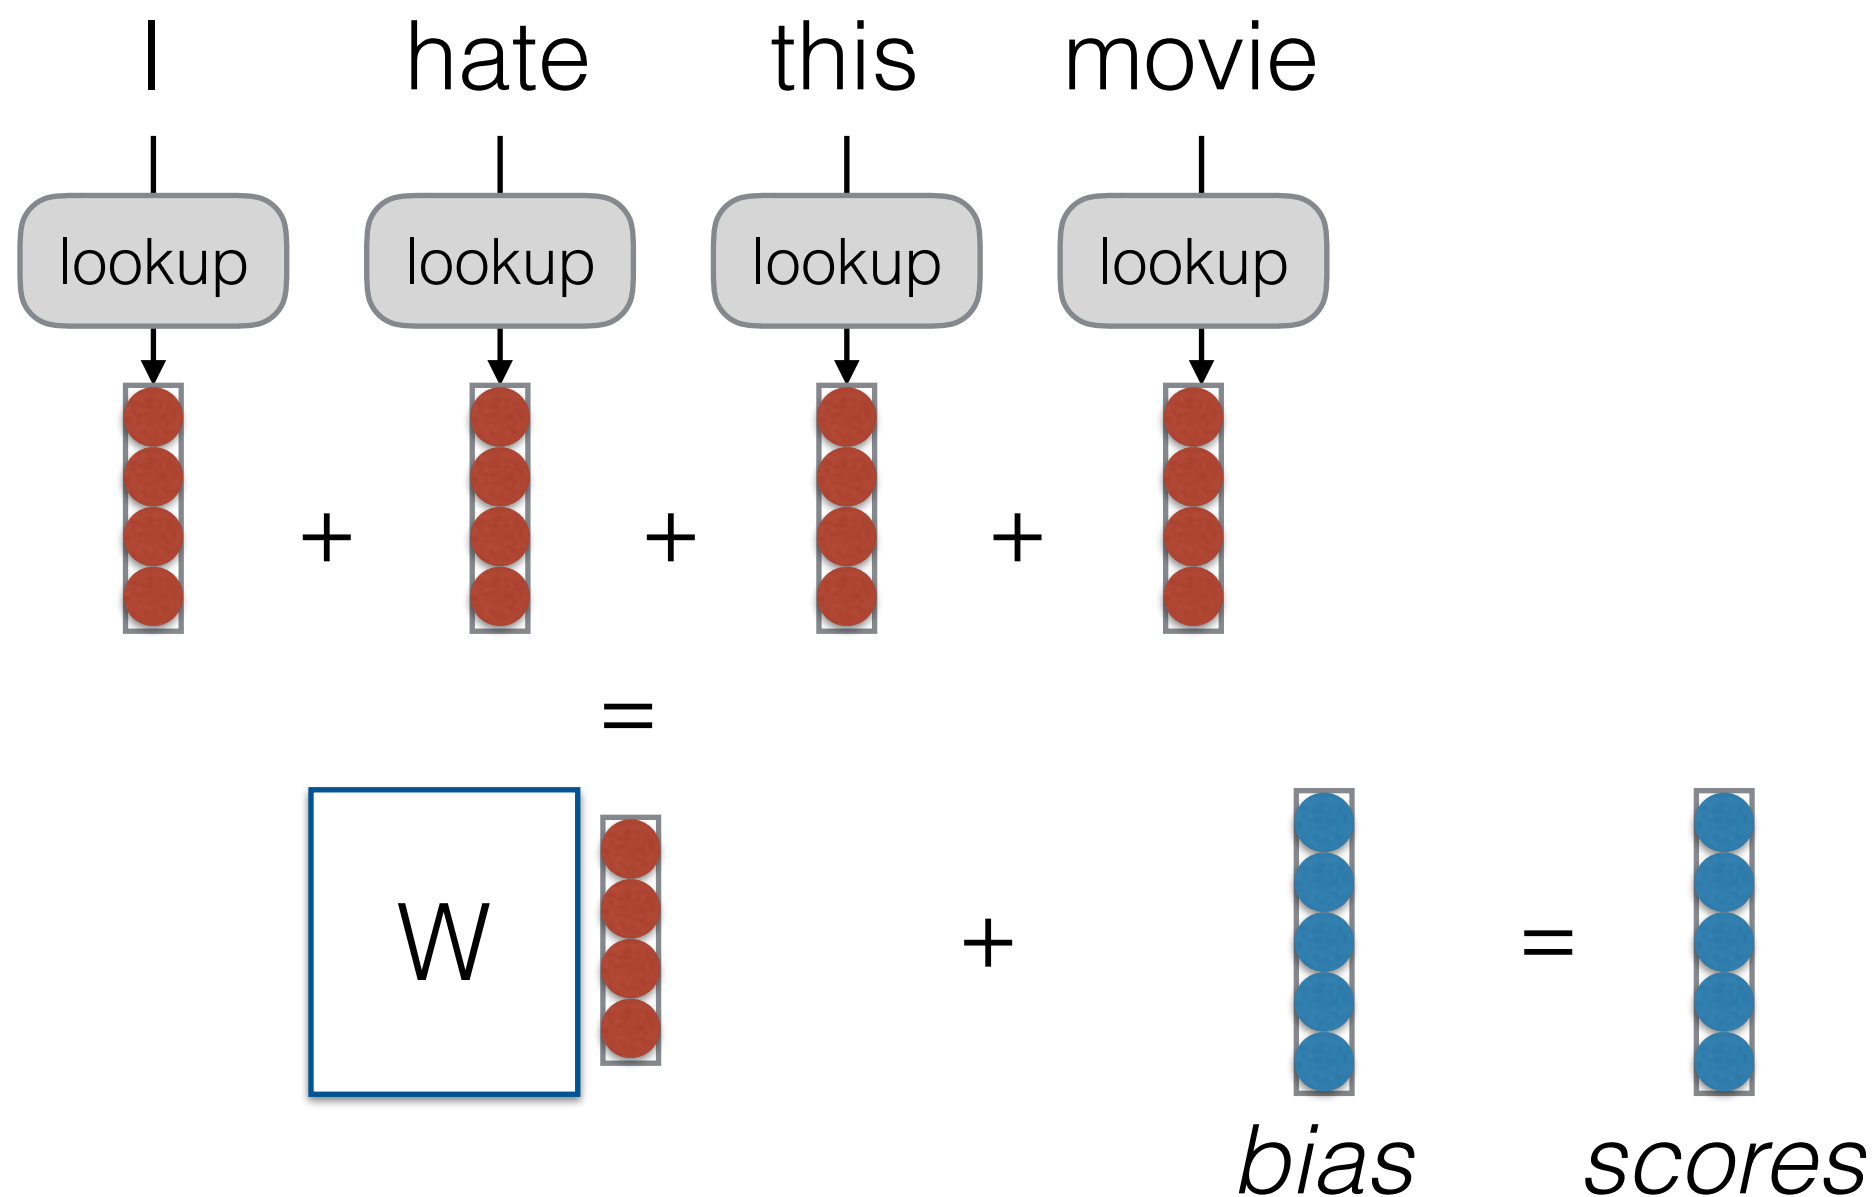
# Test Time

```
test_correct = 0.0
for words, tag in dev:
    # Define the computation graph
    scores = calc_scores(words)
    # Calculate the actual values
    score_values = scores.npvalue()
    # Find the tag with the highest score, and grade it
    predict = np.argmax(score_values)
    if predict == tag:
        test_correct += 1
print("iter %r: test acc=%.4f" % (ITER, test_correct/len(dev)))
```

Code Walk!  
(bow.ipynb)




# Continuous Bag of Words (CBOW)




Code Walk!  
(cbow.ipynb)

# Deep CBOW



Code Walk!  
(deep-cbow.ipynb)

# Bi-directional LSTM



# Builders:

## Convenience Classes for RNN, etc.

- Model definition time

```
fwdLSTM = dy.LSTMBuilder(NUM_LAYERS,  
                           EMBEDDING_SIZE,  
                           HIDDEN_SIZE,  
                           model)
```

- Training/testing time

*# Get the initial state*

```
fwd_state = fwdLSTM.initial_state()
```

*# Add the words one at a time*

```
for word_emb in word_embs:
```


```
    fwd_state = fwd_state.add_input(word_emb)
```

*# Create the output as an expression*

```
fwd_output = fwd_state.output()
```

Code Walk!  
(Istm.ipynb)

# Tree-structured RNN/LSTM





Code Walk!  
(tree-class.ipynb)

# Efficiency Tricks: Operation Batching

# Efficiency Tricks: Mini-batching

- On modern hardware 10 operations of size 1 is **much slower than** 1 operation of size 10
- Minibatching combines together smaller operations into one big one

# Minibatching

Operations w/o Minibatching

$$\tanh\left(\begin{array}{c} W \\ \begin{array}{|c|c|c|} \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \end{array} \end{array} \begin{array}{c} \mathbf{x}_1 \\ \begin{array}{|c|} \hline \bullet \\ \hline \bullet \\ \hline \bullet \\ \hline \end{array} \end{array} + \begin{array}{c} \mathbf{b} \\ \begin{array}{|c|} \hline \bullet \\ \hline \bullet \\ \hline \bullet \\ \hline \end{array} \end{array}\right) \quad \tanh\left(\begin{array}{c} W \\ \begin{array}{|c|c|c|} \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \end{array} \end{array} \begin{array}{c} \mathbf{x}_2 \\ \begin{array}{|c|} \hline \bullet \\ \hline \bullet \\ \hline \bullet \\ \hline \end{array} \end{array} + \begin{array}{c} \mathbf{b} \\ \begin{array}{|c|} \hline \bullet \\ \hline \bullet \\ \hline \bullet \\ \hline \end{array} \end{array}\right) \quad \tanh\left(\begin{array}{c} W \\ \begin{array}{|c|c|c|} \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \end{array} \end{array} \begin{array}{c} \mathbf{x}_3 \\ \begin{array}{|c|} \hline \bullet \\ \hline \bullet \\ \hline \bullet \\ \hline \end{array} \end{array} + \begin{array}{c} \mathbf{b} \\ \begin{array}{|c|} \hline \bullet \\ \hline \bullet \\ \hline \bullet \\ \hline \end{array} \end{array}\right)$$

Operations with Minibatching

$$\begin{array}{c} \mathbf{x}_1 \quad \mathbf{x}_2 \quad \mathbf{x}_3 \rightarrow \text{concat} \rightarrow \begin{array}{c} X \\ \begin{array}{|c|c|c|} \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \end{array} \end{array} \\ \begin{array}{c} W \\ \begin{array}{|c|c|c|} \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \end{array} \end{array} \end{array} + \begin{array}{c} \text{broadcast} \leftarrow \mathbf{b} \\ \begin{array}{c} B \\ \begin{array}{|c|c|c|} \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \bullet & \bullet & \bullet \\ \hline \end{array} \end{array} \end{array} \end{array}$$

# Manual Mini-batching

- DyNet has special minibatch operations for lookup and loss functions, everything else automatic
- You need to:
  - Group sentences into a mini batch (optionally, for efficiency group sentences by length)
  - Select the “t”th word in each sentence, and send them to the lookup and loss functions

# Mini-batched Code Example

---

```
1 # in_words is a tuple (word_1, word_2)
2 # out_label is an output label
3 word_1 = E[in_words[0]]
4 word_2 = E[in_words[1]]
5 scores_sym = W*dy.concatenate([word_1, word_2])+b
6 loss_sym = dy.pickneglogsoftmax(scores_sym, out_label)
```

---

(a) Non-minibatched classification.


```
1 # in_words is a list [(word_{1,1}, word_{1,2}), (word_{2,1}, word_{2,2}), ...]
2 # out_labels is a list of output labels [label_1, label_2, ...]
3 word_1_batch = dy.lookup_batch(E, [x[0] for x in in_words])
4 word_2_batch = dy.lookup_batch(E, [x[1] for x in in_words])
5 scores_sym = W*dy.concatenate([word_1_batch, word_2_batch])+b
6 loss_sym = dy.sum_batches( dy.pickneglogsoftmax_batch(scores_sym, out_labels) )
```

---


(b) Minibatched classification.

# But What about These?


## Words



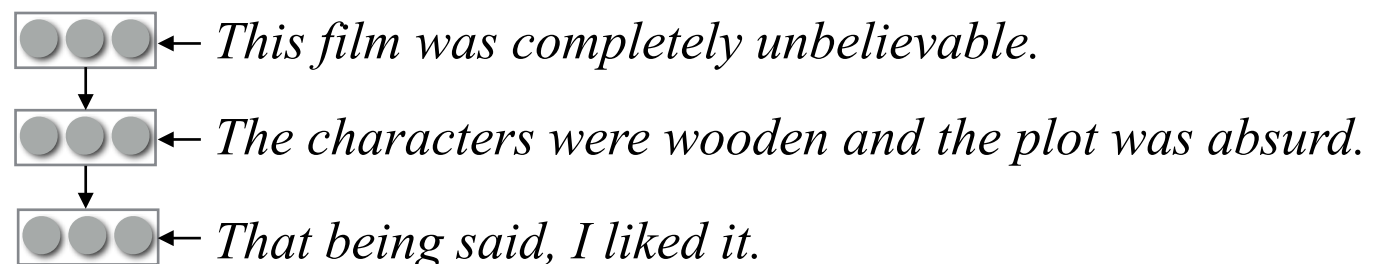
## Sentences



## Phrases



## Documents



# Automatic Mini-batching!

Three input sequences,  
different lengths.



- TensorFlow Fold (complicated combinators)
- DyNet Autobatch (basically effortless implementation)



# Autobatching Algorithm

- for each minibatch:
  - for each data point in mini-batch:
    - **define/add data**
  - **sum losses**
  - **forward** (autobatch engine does magic!)
  - **backward**
  - **update**

# Speed Improvements




Table 1: Sentences/second on various training tasks for increasingly challenging batching scenarios.

Task	CPU			GPU		
	NOAUTO	BYDEPTH	BYAGENDA	NOAUTO	BYDEPTH	BYAGENDA
BiLSTM	16.8	139	<b>156</b>	56.2	337	<b>367</b>
BiLSTM w/ char	15.7	93.8	<b>132</b>	43.2	183	<b>275</b>
TreeLSTM	50.2	348	<b>357</b>	76.5	<b>672</b>	661
Transition-Parsing	16.8	61.0	<b>61.2</b>	33.0	89.5	<b>90.1</b>

# Questions?

<https://github.com/neubig/lxmls-2017>

<https://github.com/clab/dynet>

# Supplementary Material

# Dynamic+Immediate Evaluation (PyTorch, Chainer)

- for each data point:
  - **define/add data/forward**
  - **backward**
  - **update**

# Dynamic+Lazy Evaluation (DyNet)

- for each data point:
  - **define/add data**
  - **forward**
  - **backward**
  - **update**

# Advantages/Disadvantages of Dynamic+Immediate Evaluation

- **Advantages:**

- Easy to implement nets with varying structure, API is closer to standard Python/C++

- **Disadvantages:**

- Cannot be optimized at definition time
- Harder to schedule of data transfer, etc.

# Advantages/Disadvantages of Dynamic+Lazy Evaluation

- **Advantages:**

- Easy to implement nets with varying structure
- API is closer to standard Python/C++
- Can be optimized at definition time (see end of presentation!)

- **Disadvantages:**

- Harder to schedule of data transfer, etc.